

Cenomi Mall Concierge — Competitive Position vs AI Findr

Prepared: 25 March 2026
Audience: Client / Stakeholder

How Close Are We with AI Findr?

AI Quality Dimension	Parity	Status
Conversational AI core	~87%	Cenomi leads
Retrieval quality	~85%	Cenomi leads — three-layer hybrid + scene-augmented vectors; BM25 fusion in progress
Intent classification depth	Cenomi leads	Three-tier hybrid (rule fast-path + LLM fallback); AI Findr is single-pass LLM only
Hallucination control	Cenomi leads	Post-generation validation against live canonical data; AI Findr’s grounding is unpublished
Real-time session adaptation	Cenomi leads	Mid-session behavioral tuning from feedback; AI Findr improves offline only

Where We Win

- **Multi-turn scene memory** — tracks companions, occasion, budget, and visit plan across every turn; AI Findr answers one query at a time
- **Dual-flow routing** — concierge mode for planning, factual mode for lookups; AI Findr is a single flow
- **Hallucination guard** — every LLM output validated against live mall data before delivery; AI Findr’s grounding is opaque
- **Real-time session adaptation** — feedback signals (explicit and implicit) adjust behavior mid-session; AI Findr only improves offline
- **Playbook scenario engine** — “romantic dinner” and “family outing” produce structurally different responses over the same data
- **Full data ownership** — Gulf occasions, local brands, Saudi cultural context embedded natively; AI Findr is a generic platform

- **Data residency** — all visitor data stays on Cenomi infrastructure; AI Findr requires Enterprise tier for on-premises

Where AI Findr Currently Leads

- **Visual responses** — interactive maps, store cards with images, booking CTAs
- **Arabic** — auto-detect multilingual with regional tone; we have no Arabic support yet
- **Analytics** — real-time dashboard for top queries, trends, and content gaps
- **Channels** — WhatsApp, kiosk, digital signage, voice, mobile app; we are web-only today

What AI Findr Doesn’t Have — Our Differentiators

Capability	What It Means for Visitors
Scene Memory	The concierge remembers who you’re with, your occasion, budget, and shopping task — across every turn of the conversation
Concierge vs Factual routing	Planning queries (“help me plan a date night”) and lookup queries (“where is Zara?”) are handled by separate reasoning pipelines, not a one-size-fits-all flow
Hallucination guard	Store names, hours, movies, and offers in every response are validated against live mall data before the user sees them
Real-time feedback adaptation	If a visitor corrects the bot mid-session, behavior adjusts immediately — no waiting for a next-day model update
Playbook scenario engine	Context-matched scenarios (anniversary, family day, back-to-school) change entity ranking, tone, and response structure dynamically
Scene-augmented vector search	The same query (“something for dinner”) returns different results for an anniversary couple vs a family with kids — context shapes retrieval, not just the response
Cross-mall brand search	“Is Zara in any Cenomi mall?” searches across all properties with a single validated answer
Quality evaluator	Every response is automatically scored on accuracy, relevance, and honesty in the background — without slowing delivery
Full ownership	Prompts, rules, playbooks, and entity data are Cenomi’s — Gulf/Saudi context is embedded natively, not approximated by a generic vendor

Capability	What It Means for Visitors
Data stays in-house	All visitor interactions processed on Cenomi infrastructure by default

In Progress (Roadmap)

Capability	Timeline	What It Adds
Hybrid retrieval (BM25 + vector fusion)	Pre-demo	Exact brand name queries that slip through vector search are caught by keyword index — higher recall across all query types
Re-ranking	Pre-demo	Top results for ambiguous queries reordered by true relevance — measurably better recommendations
Arabic language support	Post-demo	Full Arabic conversation with Gulf dialect handling and bilingual retrieval
Analytics dashboard	Post-demo	Business teams see top queries, trends, and content gaps without engineering involvement

Bottom Line

The Cenomi concierge is architecturally ahead of AI Findr on the AI reasoning layer — the part that determines whether visitors get genuinely useful, contextually relevant answers or generic search results. The gap to close before full parity is the **delivery layer**: channels, visual format, and Arabic support. That roadmap is defined, sequenced, and underway.